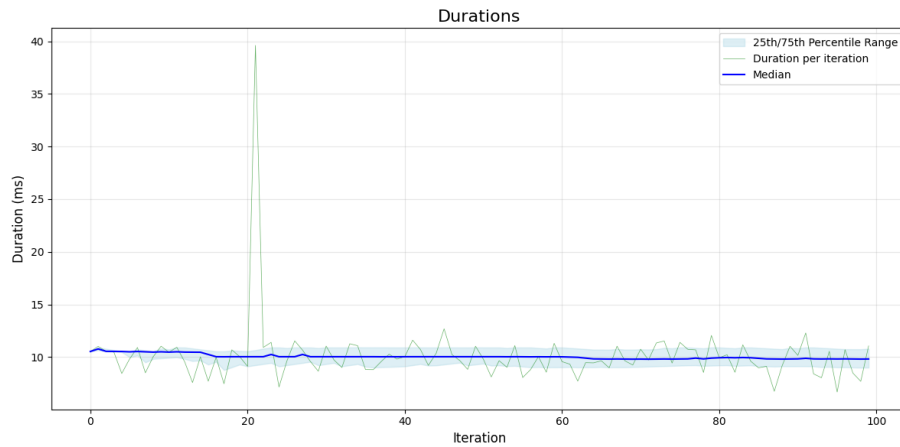


## Parallelism ablation

### Baseline

Durations for the baseline (no parallelism) implementation:



The median duration is 10 ms.

### Parallelism

We implement parallelism as follows:

1. Use the `DestroyMethod` to generate a subset of size `subset_size * parallelism`, where `parallelism` is the amount of threads/processes to use.
2. Split the `flat_subset` into `parallelism` subsets of size `subset_size`.
3. Copy `PathTable` and `instance` to each worker thread/process.
4. Run standard iteration using `LowLevelSolver`.
5. Get lowest `new_num_collisions` from worker threads.
6. If it is better than the current `num_collision`, update the solver's `path_table` and `instance` attributes.

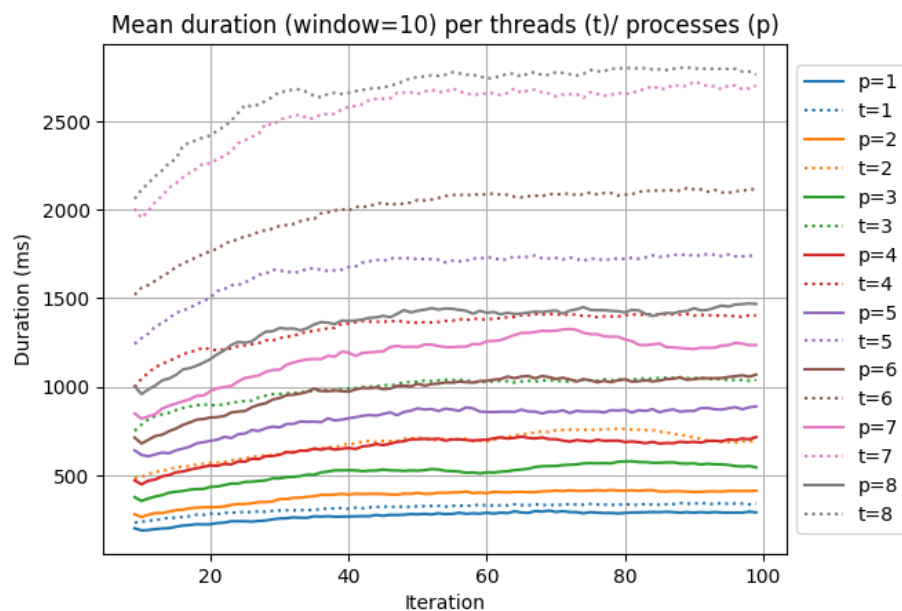
### Correctness

We tested both threads and processes by using `parallelism=1` and running the parallel solver along a standard solver. After each iteration we compared the state of both and asserted they matched.

We ran this test for 10 iterations.

## Performance

### Duration

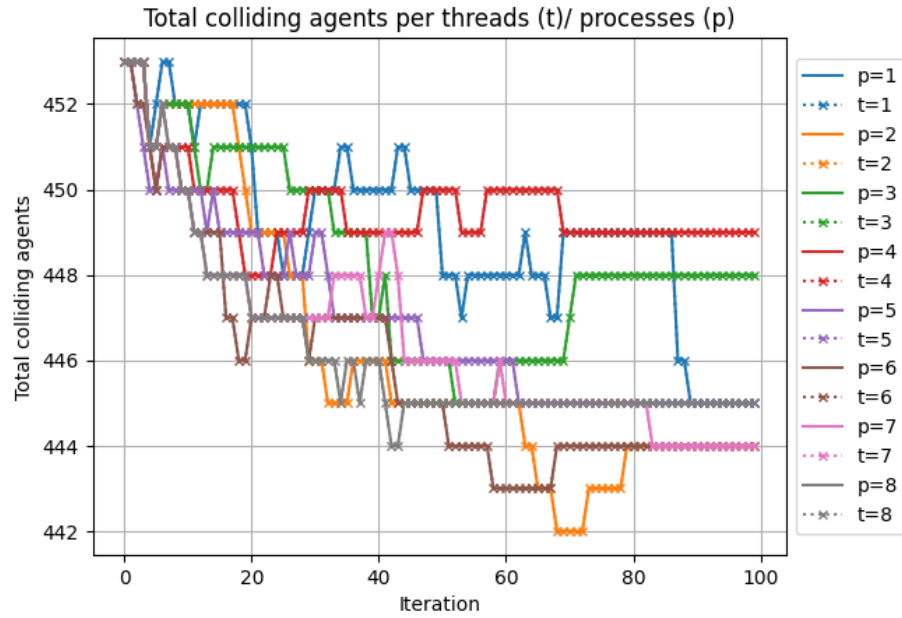


Durations per iteration are noisy, so we plot the rolling mean with a window of size 10.

Each color is for a `parallelism` amount. A solid line uses processes while dotted lines use threads.

### Collisions


We can use `parallelism=1` as a proxy for no parallelism, since their results are identical.



Y axis is the total count of agents that have at least one collision.

Solid line is for processes, dotted x is for threads. Both are identical showing implementation for both is correct.

Seeing that the unique agent collisions can go up, we changed the `num_collisions` method to calculate this value, instead of the collisions matrix sum, we get lower colliding agents:



Change in code:

```
def num_collisions_in_robots(self, num_robots = 90):
    return self.collisions_matrix.sum() // 2
    return np.sum(np.sum(self.collisions_matrix, axis=1) > 0).item()
```